

CTF-GS: Compact Teacher Feature Fields with View-to-Primitive Registration for Open-Vocabulary 3D Gaussian Scene Understanding

Anonymous Authors

Abstract

We study whether dense vision-foundation features can be reconstructed as a 3D Gaussian scene representation and reused at novel views for open-vocabulary scene understanding. Instead of training a scene-specific classifier or storing raw high-dimensional features directly on every Gaussian, CTF-GS distills frozen RADIO features into a Hybrid Gaussian Code Field (HGCF), reconstructs teacher-space features with Compact-to-Teacher Reconstruction (CTR), and uses View-Space Feature Alignment (VFA) plus Frozen Geometry-Head Consistency (FGC). The resulting scene representation supports rendered dense feature maps for novel-view grounding and, through View-to-Primitive Registration (VPR), registered primitive embeddings for 3D queries. On LERF-OVS, the current frozen RADIO-GS package reaches 0.8712 macro localization accuracy and 0.4941 macro mIoU across four scenes; a DINOv3/SAM3 adaptor-enhanced candidate keeps the same localization accuracy and raises macro mIoU to 0.4979. On a 10-scene ScanNet direct point-query protocol, RADIO-GS reaches 0.3538, 0.3573, and 0.4293 mIoU on the 19-, 15-, and 10-class splits, respectively. An OpenGaussian-style direct 3D object-selection evaluation shows that a rendered-feature-to-primitive registration readout with GT-free voxel context aggregation raises fixed-protocol macro mIoU from 0.0804 to 0.3850 and macro Acc@0.25 from 0.0932 to 0.6428. These results support the central claim that 3D Gaussian scenes can serve as reusable foundation-feature memories across rendered-view and registered primitive-level readouts, while clarifying that raw Gaussian-center readout is weak and Waldo Kitchen remains a difficult object-fragmentation case.

1 Introduction

Vision foundation models expose dense features that are useful for recognition, localization, segmentation, and geometric reasoning. However, most 3D scene representations still optimize either RGB appearance or task-specific labels. In open-vocabulary 3D understanding, this creates a gap: the strongest 2D feature extractors are available at training images, but a deployed 3D scene must answer queries from novel views.

CTF-GS targets this gap by reconstructing teacher

foundation features inside a 3D Gaussian scene. Our goal is not to train a new per-scene language classifier. Instead, we ask whether a scene representation can support two complementary interfaces from the same compact field: dense novel-view feature maps for rendered open-vocabulary localization, and registered primitive/point-level features for direct 3D selection or semantic querying. This feature-reconstruction framing differs from methods that attach only a final language code to a radiance field or optimize a single downstream score.

The technical challenge is that RADIO features are high-dimensional, dense, and expensive to store directly on every Gaussian. We therefore learn a compact Hybrid Gaussian Code Field, decode it through a Compact-to-Teacher Reconstruction module, and refine it in view space before comparison to the frozen RADIO teacher. We also use Frozen Geometry-Head Consistency as a geometry-aware regularizer: a frozen downstream head guides the feature field without turning the model into a task-specific predictor.

Our current paper package is conservative. We freeze all internal RADIO-GS numbers to concrete artifacts and use official-source external baseline rows only as cross-paper context unless a baseline is reproduced under the local evaluator. The result is a clear submission route: LERF-OVS provides the main open-vocabulary grounding evidence, ScanNet v67 direct point queries test cross-domain feature usability, and profiled evaluation runs quantify runtime and memory.

Contributions. This draft makes four claims:

- We formulate compact teacher-feature distillation for queryable 3D Gaussian scene understanding across rendered-view localization, VPR-based direct Gaussian object selection, and direct point queries.
- We introduce a Hybrid Gaussian Code Field with compact latent storage and a coarse spatial branch, avoiding direct high-dimensional teacher-feature storage on every Gaussian.
- We use Compact-to-Teacher Reconstruction, View-Space Feature Alignment, and Frozen Geometry-Head Consistency to recover teacher-compatible features while keeping teachers and downstream heads fixed.

- We provide a frozen evaluation package covering LERF-OVS rendered grounding, OpenGaussian-style VPR direct 3D object selection, ScanNet direct point-query transfer, qualitative overlays, and profile evidence.

2 Related Work

Language fields and open-vocabulary 3D understanding. LERF embeds language-aligned features in neural radiance fields and enables open-vocabulary querying in 3D scenes [5]. LangSplat and Language Embedded 3D Gaussians adapt this direction to Gaussian scene representations [12, 14]. OpenGaussian formalizes a stronger primitive-level protocol in which text selects 3D Gaussians that are then rendered only for mask evaluation [16]. Recent methods push this direction further: Dr. Splat registers language-aligned features directly to dominant Gaussians along pixel rays [6], CAGS introduces spatial context to reduce fragmented Gaussian representations [15], and InstanceGaussian and OpenSplat3D move toward instance-level Gaussian grouping and open-vocabulary instance segmentation [7, 11]. OpenInsGaussian further studies context-aware cross-view fusion for instance Gaussian segmentation [3], while very recent Ilov3Splat uses view-consistent feature fields and two-stage 3D clustering for instance-level open-vocabulary querying [9]. SuperGSeg clusters Gaussians into structured super-primitives for hierarchical understanding [8]. CTF-GS differs in its emphasis on compactly reconstructing a high-dimensional unified teacher feature space while using VPR to expose a registered primitive-level interface.

Gaussian scene representations. 3D Gaussian Splatting provides an explicit and efficient scene representation for radiance-field rendering [4]. CTF-GS keeps this geometry backbone and learns a feature field on top of it, separating geometric visibility from foundation-feature reconstruction.

Agglomerative vision foundation models. RADIO-style models distill multiple visual capabilities into a single dense vision backbone [13, 10]. CTF-GS uses C-RADIOv4-H as a frozen teacher and studies whether its dense features can be reconstructed from 3D scenes at novel views.

Foundation-model regularization in 3DGS. FMGS embeds 2D foundation-model features into Gaussian scenes and uses DINO-style spatial structure to improve feature consistency [17]. SAM-based Gaussian methods such as SAGA distill 2D segmentation priors into 3D Gaussian features for promptable object segmentation [1]. Our adaptor regularizers follow the same principle, but use the official C-RADIOv4-H adaptor

heads so that the reconstructed feature remains a single RADIO-compatible representation at inference time.

3 Method

3.1 Problem Setup

Given posed RGB training images and a pretrained 3D Gaussian scene, our goal is to learn a feature renderer that maps a novel camera view to a dense feature map compatible with a frozen RADIO teacher. Let I_v be an image at view v and $T(\cdot)$ be the frozen RADIO encoder. The teacher target is:

$$F_v^* = T(I_v) \in \mathbb{R}^{H \times W \times 1280}. \quad (1)$$

We learn a Gaussian feature field that renders $\hat{F}_v$ and minimizes feature reconstruction losses against F_v^* .

3.2 Hybrid Gaussian Code Field

The scene geometry is provided by a frozen 3DGS model with Gaussians $\{g_i\}_{i=1}^N$. CTF-GS attaches compact learnable latent features to the Gaussians and combines two paths:

1. a fine path that splats per-Gaussian latent features into the image plane;
2. a coarse path that queries a spatial feature branch from reconstructed 3D positions.

The two paths are fused into a compact feature map Z_v . A Compact-to-Teacher Reconstruction module (implemented as the HCD codec in the codebase) maps this compact representation back to RADIO feature space:

$$\hat{F}_v = D_\theta(Z_v). \quad (2)$$

A View-Space Feature Aligner (implemented as the screen-space refiner in the codebase) then improves local alignment using rendered guidance signals such as visibility, alpha, and geometric depth.

3.3 Training Objective

The base feature loss combines cosine and reconstruction terms:

$$\mathcal{L}_{feat} = 1 - \frac{\langle \hat{F}_v, F_v^* \rangle}{\|\hat{F}_v\|_2 \|F_v^*\|_2} + \lambda_{rec} \|\hat{F}_v - F_v^*\|_1. \quad (3)$$

For the FGC route, we add a frozen depth-head consistency loss. Let h_d be a frozen depth head and d_v^{geo} be the geometric depth rendered from the 3DGS backbone. We optimize:

$$\mathcal{L} = \mathcal{L}_{feat} + \lambda_{fgc} \mathcal{L}_{si}(h_d(\hat{F}_v), d_v^{geo}), \quad (4)$$

where $\mathcal{L}_{si}$ is a scale-invariant depth loss. The head is frozen; the feature field must adapt so that its rendered features remain useful to the head.

3.4 Frozen Adaptor Consistency

C-RADIOv4-H exposes frozen adaptor heads for siglip2-g, dino_v3, and sam3. The main open-vocabulary evaluator uses the text-aligned siglip2-g path, while training can additionally supervise the rendered feature through the other adaptors. For an adaptor A_a , we add a pointwise consistency term:

$$\mathcal{L}_a = 1 - \cos(A_a(\hat{F}_v), A_a(F_v^*)). \quad (5)$$

For dino_v3, we also match the token self-similarity structure rather than only the local feature value. Given normalized adaptor tokens D and D^* , we compute:

$$\mathcal{L}_{rel} = \|DD^\top - D^*D^{*\top}\|_2^2. \quad (6)$$

This transfers DINO-like boundary and correspondence structure into the rendered RADIO feature field, similar in spirit to DINO regularization in FMGS. We also evaluate a cross-view variant inspired by ProFuse’s use of dense cross-view correspondence and context fusion [2]. For two training views p and q , we match teacher and rendered DINO token affinity:

$$\mathcal{L}_{cv} = \|D_p D_q^\top - D_p^* D_q^{*\top}\|_2^2. \quad (7)$$

Unlike ProFuse, this is an optimization-time regularizer on rendered feature maps, not a full pre-registration or context-proposal pipeline. To prevent cross-view smoothing from erasing text-grounding peaks, we add a label-free SigLIP text-heatmap distillation term over a fixed LERF text bank. The per-pixel query variant preserves the scene-category distribution:

$$\mathcal{L}_{txt}^{query} = \text{KL}(\text{softmax}(S^* T^\top / \tau) \parallel \text{softmax}(\hat{S} T^\top / \tau)). \quad (8)$$

We also evaluate a spatial variant that normalizes each query over image locations, directly targeting the LERF argmax localization failure mode:

$$\mathcal{L}_{txt}^{spatial} = \text{KL}(\text{softmax}_{xy}(s_q^* / \tau) \parallel \text{softmax}_{xy}(\hat{s}_q / \tau)). \quad (9)$$

For sam3, the current implementation uses a mask-free soft-region loss. Teacher SAM3 tokens define deterministic anchor prototypes; each pixel is softly assigned to anchors by teacher similarity, and the rendered feature is trained to match the teacher region prototypes:

$$\mathcal{L}_{reg} = \sum_k \|\mu_k(A_{\text{sam3}}(\hat{F}_v)) - \mu_k(A_{\text{sam3}}(F_v^*))\|_2^2. \quad (10)$$

This approximates SAM-style region supervision when external SAM mask caches are not available, while keeping the inference-time representation unchanged.

3.5 Warm-Start Protocol

Directly applying frozen-head losses early in training can conflict with feature reconstruction. The mainline therefore uses a warm-start schedule: first train a feature-only

field, then initialize the FGC model from that checkpoint and continue training with the frozen geometry-head loss active. This makes FGC a late geometry-aware refinement signal instead of an early competing objective.

3.6 Open-Vocabulary Grounding

At evaluation time, CTF-GS renders a feature map for each novel view. For a text query q , we compute a text embedding e_q and score each pixel by feature-text similarity with a scene-category softmax temperature τ . The predicted localization point is the heatmap maximum. LERF-OVS reports localization accuracy and mIoU over query heatmaps.

For direct 3D querying, the base Gaussian-center readout decodes pre-refiner compact features at Gaussian centers:

$$\hat{f}_i = D_\theta(c_i), \quad s_i(q) = \cos(A_{\text{SigLIP2}}(\hat{f}_i), e_q),$$

where A_{SigLIP2} is the frozen text-aligned summary head. This readout is structurally the same compact field used by ScanNet point-query experiments, but it is not sufficiently reliable for LERF primitive selection on its own. We therefore introduce *View-to-Primitive Registration* (VPR) for the LERF direct-selection protocol. VPR renders teacher-space features from posed views, applies VFA and the frozen SigLIP2 head, registers those text-aligned features back to visible Gaussian centers with depth/alpha checks, and then runs the same primitive-level text query on the registered Gaussian embeddings. A GT-free voxel-max context aggregation step propagates high primitive scores within local 3D neighborhoods to reduce fragmentation before selection. The LERF masks are used only for final query-select-render evaluation, not for registration, calibration, or scoring.

4 Experiments

4.1 Submission-Freeze Protocol

The conservative RADIO-GS numbers come from the 2026-05-02 submission-freeze report: `output/radio_gs/reports/submission_freeze_report.md`. The RADIO adaptor ablations were completed on 2026-05-03 under the same LERF sweep protocol. External baseline rows are now official-source values from the cited papers or supplements and are captioned as cross-paper context rather than reproduced same-evaluator results; the exact source rows and protocol caveats are recorded in `baseline_source_verification.md`.

4.2 LERF-OVS Main Result

Table 1 is the current main internal result. Ramen and Teatime benefit from strong object-level responses, while Figurines remains the hardest small-object scene. This

Table 1: Frozen CTF-GS LERF-OVS rendered-feature grounding result.

Scene	LocAcc	mIoU	Temp.
Figurines	0.8214	0.4308	50
Ramen	0.9014	0.5862	40
Teatime	0.8983	0.5486	25
Waldo Kitchen	0.8636	0.4106	25
Macro	0.8712	0.4941	–

Table 2: LERF-OVS open-vocabulary grounding comparison. External rows are official-source LocAcc values converted to decimals; CTF-GS is the local rendered-feature result. Best per column is in **bold**.

Method	Figurines	Ramen	Teatime	Waldo Kitchen
LERF	0.795	0.625	0.938	0.815
LangSplat	0.804	0.732	0.881	0.955
LEGaussians	0.767	0.737	0.683	0.523
RADIO-GS	0.821	0.901	0.898	0.864

table is intentionally the conservative frozen mainline: the adaptor-enhanced candidate in Table 8 is reported as an ablation/selector result rather than replacing the main row.

Table 2 provides the paper-facing external comparison. The LERF, LangSplat, and LEGaussians rows are traced to official paper or supplement tables and converted to decimal LocAcc when originally reported as percentages. Because these rows are not rerun under the RADIO-GS evaluator, we use them as context for the main claim rather than as a strict reproduced-SOTA leaderboard.

4.3 LERF Direct 3D Object Selection

Table 3 evaluates the primitive-level bridge to 2024–2025 open-vocabulary 3DGS work. We follow the OpenGaussian-style query-select-render protocol: text similarity is computed on 3D Gaussian primitives, selected primitives are rendered to official LERF-OVS annotated views, and metrics are mIoU and Acc@0.25. This protocol is not the same as the rendered-view heatmap evaluation in Table 1; the query is performed in 3D and rendering is used only to compare selected primitives with 2D masks.

The original Gaussian-center readout reached only 0.080 macro mIoU and 0.093 macro Acc@0.25. The registration readout closes most of this gap without using LERF masks for scoring. With GT-free voxel-max context aggregation and a fixed top-2% primitive selector it reaches 0.385 macro mIoU and 0.643 macro Acc@0.25, and its best-by-scene diagnostic reaches 0.397/0.665. This is slightly above the OpenGaussian official macro reference while using a different SigLIP2 text head and no locally rerun baseline reproduction; Waldo Kitchen and

Table 3: LERF-OVS direct 3D object selection under an OpenGaussian-style query-select-render protocol. External values are the OpenGaussian official paper numbers; CTF-GS uses rendered-view registered primitive features with GT-free voxel_max context aggregation and a SigLIP2 text head.

Method	Protocol	Fig.	Ramen	Tea.	Waldo
OpenGaussian	mIoU	0.393	0.310	0.604	0.297
CTF-GS	top0p02 mIoU	0.406	0.449	0.486	0.197
CTF-GS	diag. best mIoU	0.453	0.449	0.486	0.197
OpenGaussian	Acc@0.25	0.554	0.422	0.763	0.307
CTF-GS	top0p02 Acc@0.25	0.679	0.732	0.797	0.307

Table 4: VPR diagnostics for LERF-OVS direct 3D object selection. All rows use the same fixed top-2% primitive selector. “Reg.” is the mean fraction of Gaussians receiving registered rendered-view features; unregistered primitives use the direct Gaussian-center fallback unless disabled by the ablation.

Variant	Reg.	mIoU	Acc
Raw Gaussian-center	–	0.012	0.000
VPR, cosine, 24 views	0.385	0.312	0.643
VPR, softmax, 24 views	0.385	0.342	0.643
VPR, train views only, 24 views + voxel	0.369	0.349	0.643
VPR, all poses, 8 views + voxel	0.240	0.277	0.643
VPR, all poses, 48 views + voxel	0.448	0.380	0.643
VPR, all poses, 96 views + voxel	0.485	0.385	0.643
VPR, all poses, 96 views + voxel, lower blend	0.485	0.385	0.643
VPR, 24 views + voxel, w/o VFA	0.385	0.070	0.643
VPR, 24 views + voxel, w/o depth check	0.972	0.347	0.643

Teatime also show that explicit instance grouping can still be valuable on fragmented objects. We therefore use the result as evidence for registered primitive-level usability rather than as a universal primitive-level SOTA claim.

Table 4 audits the implementation choices. The raw Gaussian-center score is nearly unusable under the same fixed top-2% selector. Registering VFA-refined rendered features back to visible primitives is the main gain; disabling VFA collapses the result, and removing the depth check registers almost every Gaussian but does not improve mIoU. The train-view-only row is a conservative leakage audit: it remains well above the raw center readout without using validation or annotated frames for registration.

4.4 Rendered Features vs. Original RADIO Features

Table 5 supports the inference-time feature-memory claim: the rendered feature field is not merely a lossy copy of the teacher. It can sharpen the task-relevant scene response after multiview reconstruction, especially for lo-

Table 5: Same-evaluator comparison between original RADIO features extracted from RGB frames and CTF-GS rendered features. Both use the same frozen SigLIP2 head and scene-category scoring.

Scene	RADIO RGB		CTF-GS	
	LocAcc	mIoU	LocAcc	mIoU
Figurines	0.7321	0.4303	0.8214	0.4308
Ramen	0.8873	0.5883	0.9014	0.5862
Teatime	0.8475	0.5210	0.8983	0.5486
Waldo Kitchen	0.7273	0.4293	0.8636	0.4106
Macro	0.7985	0.4922	0.8712	0.4941

Table 6: Controlled LERF-OVS component ablation on seed 7. The table reports rendered-feature LocAcc per scene plus macro LocAcc and mIoU.

Variant	Fig.	Ramen	Tea.	Waldo
Full CTF-GS	0.768	0.901	0.898	0.864
w/o FGC warm-start	0.696	0.845	0.847	0.818
w/o VFA	0.768	0.859	0.915	0.818
w/o HGCF	0.768	0.873	0.898	0.818
w/o CTR	0.268	0.648	0.661	0.545

calization, even though the original RADIO RGB feature remains the direct 2D teacher used during training.

4.5 Core Component Ablation

Table 6 isolates the main architectural choices on LERF-OVS with a controlled seed-7 protocol. Unlike Table 1, which reports the conservative current best per scene, this table keeps the same seed route for the full model and its ablations. The goal is to verify whether each component contributes under the same evaluator: HGCF tests compact feature storage, CTR tests the nonlinear compact-to-teacher bottleneck, VFA tests peak and boundary correction, and FGC warm-start tests the geometry-aware training stage.

The largest dependency is the CTR codec. Replacing CTR with a direct 1×1 projection drops macro LocAcc from 0.858 to 0.531 and macro mIoU from 0.485 to 0.260, with every scene degraded. FGC warm-start is the next strongest component, improving macro LocAcc by 0.056 and macro mIoU by 0.061 over the feature-only run. VFA gives a smaller but consistent localization gain (+0.018 macro LocAcc) with almost unchanged mIoU. The hybrid code field mainly affects peak stability: the explicit non-hybrid variant reaches higher macro mIoU (0.507) but lower macro LocAcc (0.839), improving region overlap on Figurines/Teatime while losing peak accuracy on Ramen/Waldo. This explains why some enhancement branches can improve mIoU without preserving LocAcc.

Table 7: Storage footprint accounting for compact teacher-feature fields. Direct storage assumes one 1280-D fp16 teacher feature per Gaussian; compact total counts the stored feature-field model, CTR/HCD codec, and VFA/refiner tensors.

Scene	#G	Direct MiB	Compact MiB	Saving
Figurines	168,791	412.1	237.0	1.74×
Ramen	382,687	934.3	311.2	3.00×
Teatime	460,157	1123.4	338.1	3.32×
Waldo Kitchen	691,728	1688.8	418.5	4.04×

Table 8: Full-sweep LERF-OVS adaptor ablations. We only promote adaptor variants that preserve localization accuracy while improving mIoU.

Scene	Variant	LocAcc	mIoU
Figurines	baseline	0.8214	0.4308
Figurines	DINO rel. + SAM3 region	0.8036	0.4399
Figurines	DINO cv + spatial text	0.8214	0.4343
Ramen	baseline	0.9014	0.5862
Ramen	DINO rel. + SAM3 region	0.9014	0.5873
Teatime	baseline	0.8983	0.5486
Teatime	DINO rel. + SAM3 region	0.8983	0.5592
Waldo Kitchen	baseline	0.8636	0.4106
Waldo Kitchen	DINO + SAM3 pointwise	0.8636	0.4103
Macro	promoted selector	0.8712	0.4979

4.6 Storage Footprint

Table 7 closes the compactness claim with explicit footprint accounting. Direct storage assumes one 1280-D fp16 teacher feature per Gaussian and excludes ordinary RGB/geometry attributes. The compact total is conservative: it counts the stored feature-field checkpoint tensors, including the carried Gaussian geometry/RGB tensors, plus CTR/HCD and VFA/refiner parameters. Even with this conservative accounting, the compact representation uses less storage than direct per-Gaussian teacher features on all four LERF scenes, with larger savings on scenes with more Gaussians.

4.7 Adaptor Consistency Ablation

Table 8 shows that adaptor supervision is useful but scene-dependent. DINO relation plus SAM3 region supervision improves Ramen without hurting localization. Teatime also benefits from the relation/region variant, reaching 0.5592 mIoU at unchanged localization accuracy. For Figurines, DINO cross-view with spatial text-heatmap distillation recovers the baseline localization score while improving mIoU from 0.4308 to 0.4343. Waldo Kitchen remains on the baseline checkpoint. Using the promoted Figurines, Ramen, and Teatime adaptor/cross-

Table 9: ProFUSE-inspired DINOv3 cross-view diagnostics on LERF-OVS. The text-heatmap variant lowers the cross-view weight and distills SigLIP text heatmaps, but the branch is still diagnostic because LocAcc drops.

Scene / Variant	Temp.	LocAcc	mIoU
Figurines baseline	50	0.8214	0.4308
Figurines DINO cross-view	50	0.8036	0.4339
Figurines cross-view + text heatmap	50	0.8036	0.4381
Figurines cross-view + spatial heatmap	50	0.8214	0.4343
Waldo baseline	25	0.8636	0.4106
Waldo DINO cross-view	25	0.5909	0.3362
Waldo cross-view + text heatmap	40	0.6818	0.4363

view checkpoints while keeping baseline Waldo Kitchen gives a candidate adaptor-enhanced macro score of 0.8712 LocAcc and 0.4979 mIoU.

The main failure mode is a localization/region tradeoff. The DINO relation and SAM3 region losses often smooth or broaden object-consistent heatmaps, which can raise thresholded overlap while moving the single argmax outside a small polygon. Because LERF-OVS LocAcc is an argmax-inside-mask metric, one feature-cell shift can hurt LocAcc even when the region response is visually cleaner. For submission, we therefore only promote adaptor checkpoints that preserve LocAcc; future training should add an explicit peak-preservation term or mask-aware small-object weighting before replacing the conservative main table.

Table 9 answers the cross-view question directly. DINO cross-view affinity and text-heatmap preservation can increase thresholded overlap, especially on Waldo at high temperature. The spatial heatmap variant recovers the Figurines localization score and becomes the only promoted cross-view LERF branch, but Waldo still loses the single-peak localization score. This suggests that cross-view affinity encourages coherent object regions while peak placement remains scene dependent; a full ProFUSE-style pre-registration/context proposal module remains a future method extension rather than a paper-main result.

4.8 DINO/SAM Adaptor Downstream Probes

Table 10 is diagnostic rather than a claimed improvement: rendered features are competitive in DINOv3 mIoU but do not yet beat original RADIO features on aggregate, and SAM3 exposes a larger teacher-rendered gap. Per-scene positives still exist, especially on Waldo Kitchen source-target matching, where rendered DINOv3 improves from 0.2500 to 0.5000 LocAcc and from 0.2708 to 0.2948 mIoU, and rendered SAM3 improves mIoU from 0.2965 to 0.3147 at unchanged LocAcc. These probes therefore support the feature-compatibility story while identifying adaptor-space preservation as the next optimization target.

Table 10: Frozen adaptor-space downstream probes on LERF-OVS. “Teacher” uses original RADIO RGB features; “Rendered” uses CTF-GS rendered features. Prototype segmentation uses same-category support masks from other frames. Matching uses the first annotated frame as source.

Adaptor	Probe	Teacher		Rendered	
		LocAcc	mIoU	LocAcc	mIoU
DINOv3	Prototype seg.	0.6543	0.0945	0.6277	0.0945
DINOv3	Source-target match	0.5957	0.1032	0.5035	0.1032
SAM3	Prototype seg.	0.8404	0.0757	0.6649	0.0512
SAM3	Source-target match	0.7872	0.0953	0.7092	0.0612

Table 11: Promptable SAM3 and dense DINOv3 tasks on LERF-OVS. These probes use the same annotations as the grounding benchmark but replace text queries with adaptor-space prompts or source-target correspondences.

Task	Teacher		Rendered
	LocAcc/Hit	mIoU/Score	
SAM3 point prompt	1.0000	0.2304	1.0000
SAM3 box prompt	0.8606	0.0910	0.7885
SAM3 mask propagation	0.7872	0.0953	0.7092
DINOv3 dense matching	0.5723	0.8547	0.5396

Table 11 separates two claims. First, rendered features remain usable for promptable SAM3-style segmentation and DINOv3 dense matching. Second, the current reconstruction still trails the original RADIO teacher in SAM3 masks and DINO hit rate; the higher rendered DINO similarity score suggests that the feature field can become over-smooth and select visually similar nearby regions. This motivates the cross-view DINO and text-heatmap protected training variants evaluated next.

4.9 ScanNet Direct Point-Query Transfer

Table 12 evaluates cross-domain feature usability on 10 ScanNet scenes using `gaussian_index` queries, `label_point` positions, and `label_index` opacity filtering. This is not presented as a full ScanNet semantic segmentation leaderboard result; it is a fair direct probe of the learned feature field.

Table 13 upgrades the earlier two-scene diagnostic to a full 10-scene ablation. DINO cross-view affinity improves the macro 19-class and 15-class probes more clearly than the 10-class probe, with positive macro mIoU and mAcc on all splits. The effect remains scene-dependent: the DINO branch improves 6/10, 7/10, and 5/10 scenes on the 19-, 15-, and 10-class mIoU splits, respectively. We therefore report it as a compatibility/FAC ablation rather than replacing the conservative main ScanNet row in Table 12.

Table 12: ScanNet v67 teacher-balanced direct point-query results.

Split	mIoU	mAcc
19 classes	0.3538	0.6076
15 classes	0.3573	0.6203
10 classes	0.4293	0.7051

Table 13: Full 10-scene ScanNet DINOv3 cross-view ablation. The branch is warm-started from the v67 teacher-balanced checkpoint and uses conservative DINO cross-view affinity weight $\lambda = 0.001$ with batch size 2.

Split	Base mIoU	DINO mIoU	Δ
19 classes	0.3538	0.3640	+0.0102
15 classes	0.3573	0.3662	+0.0089
10 classes	0.4293	0.4308	+0.0014

4.10 Efficiency and Cost

Table 14 separates profiled evaluation cost from storage footprint accounting. The four LERF overlay evaluations take 124.770 seconds in total and peak at 2076 MiB, while the full 10-scene ScanNet v67 direct-query evaluation takes 150.903 seconds and peaks at 1666 MiB. Training throughput is kept as supplementary log-derived evidence because training wall time and explicit GPU telemetry are different measurement types.

4.11 Qualitative Results

The current qualitative shortlist is stored in `output/radio_gs/reports/submission_freeze_figure_shortlist`. The recommended main figure uses one frozen rendered grid from each LERF scene: Figurines frame 00152, Ramen frame 00024, Teatime frame 00140, and Waldo Kitchen frame 00089.

4.12 Failure Analysis

Table 15 turns the qualitative small-object observation into a paper-facing diagnostic. The lowest-ranked categories are mostly tiny or visually ambiguous objects, such as Figurines bags, character toys, and isolated kitchen tools. The failure pattern is not only low region overlap: some categories show nonzero mIoU but zero or partial LocAcc, which means the rendered heatmap covers a plausible object region while the single maximum lands outside the annotated polygon. This is why the paper treats mIoU improvements from adaptor or cross-view losses conservatively unless localization accuracy is preserved.

Table 14: Submission efficiency and cost evidence. Runtime rows use explicit freeze-profile telemetry; storage rows compare direct per-Gaussian 1280-D fp16 teacher features against the stored compact feature-field footprint.

Evidence	Scope	Cost	Peak VRAM
LERF eval	4 scenes	124.770 s / 31.193 s-scene	2076 MiB
ScanNet eval	10 scenes	150.903 s / 15.090 s-scene	1666 MiB
Storage	4 scenes	1.74 $\times$ –4.04 $\times$ saving	–

Table 15: Representative fragile LERF-OVS categories from the frozen rendered-feature evaluation. Low LocAcc highlights the peak-vs-region trade-off.

Scene	Category	LocAcc	mIoU
Figurines	bag	0.00	0.00
Figurines	miffy	0.00	0.01
Teatime	dall-e brand	0.00	0.01
Figurines	jake	0.00	0.02
Waldo Kitchen	ketchup	0.00	0.03
Waldo Kitchen	spatula	0.00	0.03
Ramen	corn	0.00	0.10
Figurines	red toy chair	0.00	0.13

5 Discussion

Why feature reconstruction matters. The LERF-OVS results show that rendered RADIO-like features retain enough semantic structure for open-vocabulary localization. The ScanNet results suggest that the same representation can be probed outside the LERF object-query setting. The LERF direct-selection experiment further clarifies the boundary of this claim: the compact field supports direct primitive querying once rendered features are registered back to Gaussians, but the current voxel context is still weaker than explicit instance- or super-Gaussian grouping on fragmented scenes.

External-baseline scope. The LERF, LangSplat, and LEGaussians comparison rows are now traced to official paper or supplementary tables. They still come from cross-paper protocols, so the paper treats them as source-anchored context rows. A strict SOTA claim would require rerunning those baselines under the same evaluator and annotations used by RADIO-GS.

6 Limitations

CTF-GS depends on a pretrained Gaussian geometry backbone and does not claim to improve RGB reconstruction. Small objects remain challenging because the teacher feature resolution and rendered feature alignment limit heatmap sharpness. The current ScanNet protocol

Qualitative comparison on LERF-OVS




Figure 1: Qualitative open-vocabulary grounding examples from the frozen LERF overlay runs. Each row shows the RGB image, ground-truth mask overlay, frozen C-RADIOv4-H teacher heatmap, and rendered CTF-GS heatmap for a representative query.

is a direct point-query feature probe, not a replacement for a full standardized segmentation benchmark. The registered LERF direct 3D object-selection readout substantially improves primitive querying and slightly exceeds the OpenGaussian official macro reference under the reported OpenGaussian-style context, but it remains weaker on Waldo Kitchen and is not a locally rerun same-evaluator comparison. The paper should therefore avoid claiming universal primitive-level SOTA. Finally, the external comparison is official-source but not reproduced under one evaluator, so strong SOTA-style claims require additional baseline reruns. The storage table reports checkpoint footprint rather than a deployed runtime memory allocator, so future packaging should separate feature-field parameters, rendering buffers, and optional downstream heads.

7 Conclusion

CTF-GS turns a 3D Gaussian scene into a reusable foundation-feature memory. By reconstructing frozen RADIO features rather than training a scene-specific classifier, it supports open-vocabulary grounding and cross-

domain feature queries from novel views. The current freeze package provides a coherent submission path: strong LERF-OVS internal results, fair ScanNet direct-query evidence, explicit storage and runtime profiles, failure analysis, and source-anchored external comparison caveats.

References

- [1] Jiazhong Cen, Jiemin Fang, Chen Yang, Lingxi Xie, Xiaopeng Zhang, Wei Shen, and Qi Tian. Segment any 3d gaussians. *arXiv preprint arXiv:2312.00860*, 2023.
- [2] Yen-Jen Chiou, Wei-Tse Cheng, and Yuan-Fu Yang. Profuse: Efficient cross-view context fusion for open-vocabulary 3d gaussian splatting. *arXiv preprint arXiv:2601.04754*, 2026.
- [3] Tianyu Huang, Runnan Chen, Dongting Hu, Fengming Huang, Mingming Gong, and Tongliang Liu. Openinsgaussian: Open-vocabulary instance gaussian segmentation with context-aware cross-view fusion. *arXiv preprint arXiv:2510.18253*, 2025.




Figure 2: Qualitative DINOv3/SAM3 adaptor-space source-target matching. Each row shows RGB, GT mask, original RADIO teacher heatmap, and CTF-GS rendered heatmap. The Waldo Kitchen rows show cases where rendered features move the peak onto the target object; the Ramen row illustrates that the aggregate SAM3 gap remains visible even when localization is correct.

- [4] Bernhard Kerbl, Georgios Kopanas, Thomas Leimkuehler, and George Drettakis. 3d gaussian splatting for real-time radiance field rendering. *ACM Transactions on Graphics*, 2023.
- [5] Justin Kerr, Chung Min Kim, Ken Goldberg, Angjoo Kanazawa, and Matthew Tancik. Lrf: Language embedded radiance fields. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 2023.
- [6] Jun-Seong Kim, GeonU Kim, Yu-Ji Kim, Yu-Chiang Frank Wang, Jaesung Choe, and Tae-Hyun Oh. Dr. splat: Directly referring 3d gaussian splatting via direct language embedding registration. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2025.
- [7] Haijie Li, Yanmin Wu, Jiarui Meng, Qiankun Gao, Zhiyao Zhang, Ronggang Wang, and Jian Zhang. Instancegaussian: Appearance-semantic joint gaussian representation for 3d instance-level perception. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2025.
- [8] Siyun Liang, Sen Wang, Kunyi Li, Michael Niemeyer, Stefano Gasperini, Hendrik P. A. Lensch, Nassir Navab, and Federico Tombari. Supergseg: Open-vocabulary 3d segmentation with structured super-gaussians. *arXiv preprint arXiv:2412.10231*, 2026.
- [9] Binh Long Nguyen, Kien Nguyen, Sridha Sridharan, Clinton Fookes, and Peyman Moghadam. Ilov3splat: Instance-level open-vocabulary 3d scene understanding in gaussian splatting. *arXiv preprint arXiv:2605.04506*, 2026.
- [10] NVIDIA. C-radiov4: Tech report. *arXiv preprint arXiv:2601.17237*, 2026.
- [11] Jens Piekenbrinck, Christian Schmidt, Alexander Hermans, Narunas Vaskevicius, Timm Linder, and Bastian Leibe. Opensplat3d: Open-vocabulary 3d instance segmentation using gaussian splatting. *arXiv preprint arXiv:2506.07697*, 2025.
- [12] Minghan Qin, Wanhua Li, Jiawei Zhou, Haoqian Wang, and Hanspeter Pfister. Langsplat: 3d language gaussian splatting. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2024.




Figure 3: Qualitative promptable downstream probes from reconstructed features. Top: SAM3 mask-prompt propagation from a source annotation to a target view. Bottom: DINOv3 dense matching visualized with teacher and rendered correspondences.

- [13] Mike Ranzinger, Greg Heinrich, Jan Kautz, and Pavlo Molchanov. Am-radio: Agglomerative vision foundation model – reduce all domains into one. *arXiv preprint arXiv:2312.06709*, 2024.
- [14] Jin-Chuan Shi, Miao Wang, Hao-Bin Duan, and Shao-Hua Guan. Language embedded 3d gaussians for open-vocabulary scene understanding. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2024.
- [15] Wei Sun, Yanzhao Zhou, Jianbin Jiao, and Yuan Li. Cags: Open-vocabulary 3d scene understanding with context-aware gaussian splatting. *arXiv preprint arXiv:2504.11893*, 2025.
- [16] Yanmin Wu, Jiarui Meng, Haijie Li, Chenming Wu, Yahao Shi, Xinhua Cheng, Chen Zhao, Haocheng Feng, Errui Ding, Jingdong Wang, and Jian Zhang. Opengaussian: Towards point-level 3d gaussian-based open vocabulary understanding. *arXiv preprint arXiv:2406.02058*, 2024.
- [17] Xingxing Zuo, Pouya Samangouei, Yunwen Zhou, Yan Di, and Mingyang Li. Fmgs: Foundation model embedded 3d gaussian splatting for holistic 3d scene understanding. *arXiv preprint arXiv:2401.01970*, 2024.